

Optimizing Breadth-First Search on the NVIDIA DGX Spark

Student: Fabio Missagia

Supervisors: Flavio Vella, Thomas Pasquali, Jacopo Raffi

Department of Information Engineering and Computer Science, University of Trento



UNIVERSITÀ
DI TRENTO

 HICREST

NVIDIA DGX Spark: general properties

Property	Value
System	NVIDIA DGX Spark
CPU	20 cores: 10× Cortex-X925 + 10× Cortex-A725 (aarch64)
GPU	NVIDIA GB10 — compute capability 12.1 (sm_121)
Memory	128 GB unified LPDDR5x, CPU–GPU coherent, ~ 273 GB/s <i>shared</i>
L2 cache	24 MiB
SMs	48
GPCs	4 (12 SMs each) — <i>measured, see next slides</i>
Max grid	$2\,147\,483\,647 \times 65\,535 \times 65\,535$

One *physical* memory pool, hardware coherent: a kernel dereferences a plain `malloc()` pointer with no `cudaMemcpy`. ~ 273 GB/s is an order of magnitude below datacenter HBM, like H100 (~ 2.0 TB/s).

NVIDIA DGX Spark: per-SM properties


Property	Value
Max threads / block	1024
Max threads / SM	1536 (48 warps)
Max blocks / SM	24
Registers / SM	65 536
Unified L1 + shared memory	128 KiB
Max shared memory / SM	102 400 B (100 KiB)
Default shared memory / block	49 152 B (48 KiB)
Max shared memory / block (opt-in)	101 376 B (99 KiB)
Reserved shared memory / block	1 024 B

L1 and shared memory are the *same* 128 KiB SRAM, split by a configurable carveout: shared memory is bought at the cost of L1.

Thread block clusters (compute capability 9.0+)

An optional level between block and grid:

- **adjacent thread blocks** are grouped into clusters
- guaranteed to be **co-scheduled simultaneously inside one GPC**
- any thread block can address **any peer's shared memory** — loads, stores and *atomics* ride an **SM-to-SM** interconnect, not L2/DRAM $\Rightarrow$ **DSMEM**
- **co-residency** is guaranteed
- hardware barrier across blocks: `cluster.sync()`



Thread block clusters on the DGX Spark

Measurement	portable	with opt-in
Cluster Size	8	12

- **8** is the maximum CUDA guarantees *portably* on any cluster-capable GPU — code using it runs anywhere.
- **12** is *this hardware's* real maximum, unlocked per kernel with `cudaFuncAttributeNonPortableClusterSizeAllowed` — at the price of code that is no longer portable. Above 12: *cluster misconfiguration*.
- `cudaOccupancyMaxPotentialClusterSize` is **per kernel**, not a device constant: the same kernel answers 8 before the opt-in and 12 after.
- 12 is *usable*, not merely launchable: DSMEM verified across all 12 ranks with a ring test (rank r reads rank $r+1 \bmod n$).

How many GPCs?

Problem. No CUDA attribute exposes the GPC count or the SM-to-GPC mapping.

Idea. Exploit the one guarantee available: *all blocks of a cluster are co-scheduled in a single GPC*. So the SMs inside one cluster are provably all in the *same* GPC.

Method. Launch of 4 clusters $\times$ 12 blocks.

- every block reads the `%smid` register and reports the **SM it ran on** and the **cluster it belongs to**
- **52 KB of shared memory per block** forces one block per SM, so a 12-block cluster covers 12 *distinct* SMs
- cluster size 12 = one GPC's SM count, so **4 clusters cover all 48 SMs** exactly once and each cluster is a GPC
- **tried many trials**, but the result is always the same and all 48 SMs are observed every time, so one launch is enough

GPC topology: 4 GPCs $\times$ 12 SMs

Example:

GPC	SM ids											
GPC 0	0	1	8	9	16	17	24	25	32	33	40	41
GPC 1	2	3	10	11	18	19	26	27	34	35	42	43
GPC 2	4	5	12	13	20	21	28	29	36	37	44	45
GPC 3	6	7	14	15	22	23	30	31	38	39	46	47

Derived and verified:

- $TPC = smid/2$, $GPC = (smid/2) \% 4$
- 24 TPCs of 2 SMs, 6 TPCs per GPC, round-robin across GPCs


Three independent cross-checks:

- 4 clusters of 12 *distinct* SMs, 48 SMs observed, none twice
- $4 \times 12 = 48$ matches `multiProcessorCount`
- the largest launchable cluster (12) *equals* one GPC's SM count

How do we measure latency?

A GPU is built to hide latency: a load does not stall the thread, this is called Memory-Level Parallelism.

The fix: make each load depend on the result of the previous one, so that the thread cannot issue the next load until the previous one has completed. This is called a **dependent pointer chase**, or pChase.



`idx = buf[idx];` the value loaded *is* the next address



In addition:

- **Sattolo cycle:** `buf` is one random cycle over every element: no short loops, nothing to prefetch.
- **steps:** how many dependent loads the timed loop performs. If far below the cycle length: no address is revisited.
- **reps:** The whole kernel is launched many times — a drift across reps exposes a cold cache.
- **Warm-up:** L1: one untimed lap to warm the cache; L2: a coalesced sweep *after* the kernel is loaded (lazy load flushes L2).

Latency Benchmark – Random pChase

Dependent random pointer chase, 16 384 steps, one thread, median of 101 reps.


Placement	GB10 (cy)	ns	H800 (cy)*
L1 Cache	43.2	17.9	32.0
SMEM	32.9	13.7	29.0
SMEM, via DSMEM API	36.9	15.3	33.0
DSMEM	180-211	75-88	181-213
L2 Cache	340.1	141.7	264–502
DRAM	965.8	401.6	656

* Luo et al., arXiv:2501.12084 — Hopper H800, same p-chase method

Why does SMEM via DSMEM cost 4 more cycles? `map_shared_rank` returns a generic pointer, so the load compiles to LD (address space resolved by hardware on every access) instead of LDS (a plain shared-memory load).


Latency Benchmark – Random pChase

Dependent random pointer chase, 16 384 steps, one thread, median of 101 reps.



Latency Benchmark – Random pChase

Same chase, but thread 0 of rank 0 accesses a *peer* block's DSMEM at rank distance d .



- **A genuine SM-to-SM path.** A remote access costs about *half* an L2 hit and a *fifth* of DRAM.
- **Flat across the cluster.** 210.9 cy at distance 1, identically at cluster sizes 2, 4, 8 and 12. Not monotone in distance: ranks pair up $\{2, 3\}, \{4, 5\}, \dots$ — the two SMs of one TPC.
- **Independently confirmed.** H800 agrees within a few cycles; since GB10's L2 and DRAM are *slower* than Hopper's while its fabric is not, DSMEM's advantage is *larger* here.

Distance does not matter


Rank 0 → its own TPC sibling (rank 1) is 210.9 cy, the *slowest* pair of all...

Rank 0 → the far rank 10 is 179.9 cy, the *fastest* pair of all...

Let's measure every pair!

Facts from the matrix

- perfectly **symmetric**
- TPC siblings have **identical rows**
- the diagonal is **not flat** (211 vs 149)
- **same results** in all 4 GPCs



If distance does not matter, what does?

Over all 36 TPC pairs the following holds with zero error:

$$\text{latency}(k \rightarrow d) = 149 + w_k + w_d$$

Maybe a **shared meeting point**?

Every request may go $\text{SM} \rightarrow \text{hub} \rightarrow \text{SM}$.

TPC	5	2	4	1	3	0
w (cy)	0	8	12	19	24	31

Independent confirmation: reading $L2$ from every SM shows the same pattern!

median L2 latency (cy), per TPC





TPC pos.	0	1	2	3	4	5
GPC 0	343	331	320	336	324	312
GPC 1	318	306	295	311	299	287
GPC 2	318	306	295	311	299	287
GPC 3	341	329	318	334	322	310

row minus its minimum = per-TPC distance

TPC pos.	0	1	2	3	4	5
GPC 0	31	19	8	24	12	0
GPC 1	31	19	8	24	12	0
GPC 2	31	19	8	24	12	0
GPC 3	31	19	8	24	12	0

Latency Benchmark – Linear pChase

Same chase, but element i points to $i + s$, $s = 4\text{ B} \dots 4\text{ KiB}$. Dashed = random chase (previous slide).



- **SMEM/DSMEM are pattern-blind**
- **L2/DRAM see only the 32 B sector**
 $k = 32/s$ loads share one L1 fill,
 $(T + (k - 1) \cdot 43.2)/k$ gives
80/118/192 cy (L2) and
155/266/490 cy (DRAM) — as
measured.
- **No hardware prefetcher:** from 32 B
on, linear = random.

Clusters are not free

The moment a launch carries a cluster dimension, the device holds **2 blocks per SM at most (96 in total)**, whatever the block size.

blockDim	total blocks (unclustered)	total blocks (clustered)	penalty	occupancy (clustered)
64	1152	96	12×	8%
128	576	96	6×	17%
256	288	96	3×	33%
512	144	96	1.5×	67%
1024	48	48	none	67%

- **Confirmed two independent ways:** `cudaOccupancyMaxActiveClusters`, and a test in which every block must observe all the others before a timeout.
- **Didn't found anything about this on the web...**

Clusters are not free

The cap also moves with the *shared memory* a block requests. Past ~ 50 KiB only one block fits per SM, so the device holds 48.

SMEM / block	cs = 1	cs = 2	cs = 4	cs = 8	cs = 12
0	96	96	96	96	96
32 KiB	96	96	96	96	96
50 KiB	48	48	48	32	48
99 KiB	48	48	48	32	48


total resident blocks device-wide (48 SMs), cs = cluster size

- A cluster never spans a GPC, and a GPC has 12 SMs. One cluster of 8 leaves 4 SMs free — too few for a second — so each of the 4 GPCs runs one cluster and strands 4 SMs: $4 \times 8 = 32$, a third of the GPU idle.
- **12 and 6 tile the machine exactly**; 8, the *portable* maximum, does not.


Throughput Benchmark

So far: *one* thread reading *one* peer. Now: *many* threads, reading the *same* peer or *different* peers.

all vs 1 (broadcast)




all vs all (ring)





- One cluster, one block per SM, 256 threads per block.
- Two access patterns, both dependent chases (one load in flight per thread):
 - pchase: each lane starts from a different random position in the same cycle, resulting in approximately 32 independent dependent-load chains per warp.
 - coalesced: lane l starts at $base + l \Rightarrow$ a warp reads 32 consecutive words, one 128 B line, a *single* request.

Latency Benchmark – Many Threads



Spreading the traffic: $6.3\times$, but only under load



GB/s	1 thr	128 thr
all vs 1	0.50	0.71
all vs all	0.55	4.45

- At 1 thread/SM the two are the same (0.50 vs 0.55): 12 threads never fill the pipe.
- Loaded, spreading is worth $6.3\times$: threads buy all vs 1 only $1.4\times$, all vs all $8.1\times$.
- Two ceilings: 0.71 GB/s per owner, 4.45 GB/s fabric-wide — already reached at 6 owners.

What's next?

1. Characterization of DSMEM

- Performance characterization and impact of DSMEM
- Limitations and assumptions of the DSMEM programming paradigm
- Theoretical upper bounds on performance improvements
- Applicability: algorithms that can and cannot exploit DSMEM

2. State-of-the-art GPU-accelerated BFS

- Microbenchmark access pattern of a BFS?
- Graph partitioning strategies
- Work-stealing techniques
- Data types and data structures for high-performance BFS

3. BFS implementation based on DSMEM

- Algorithm design informed by (1) and (2)
- Graph partitioning (probably 2D) and work stealing
- Integration of *mxint* (micro-scaling integer type), if worthwhile
- Performance evaluation; limitations and future directions